



İZMİR
KÂTİP ÇELEBİ
ÜNİVERSİTESİ

2010

EEE 450- Introduction to Machine Learning Second Application Assignment

DEPARTMENT OF ELECTRICAL AND ELECTRONICS ENGINEERING,
İZMİR KATIP ÇELEBİ UNIVERSITY

Turhan Can Kargin | 150403005 | Deadline:12-12-2019

Assignment#1:

Spiral dataset in the Weka tool

- Classify the Spiral dataset in Weka using the Linear (functions/Liblinear), Neural Network (MultiLayerPerceptron) and Support Vector Machines (functions/LibSVM) algorithms.
- The Spiral dataset can be downloaded from the Datasets page.
- Why does the Linear classifier has much lower accuracy than Neural Network and Support Vector Machine?
- Note that Weka automatically tries to determine the size of the hidden layer. It often happens that it uses too few hidden units to be able to accurately learn the concept. Try to change the hiddenLayers field from a to 72.

Solution:

LibLinear

The screenshot shows the Weka Classifier window with the LibLinear classifier selected. The 'Test options' section on the left shows 'Cross-validation' with 'Folds' set to 10. The 'Classifier output' section on the right displays the results of the stratified cross-validation.

Classifier output:

```
Time taken to build model: 0.04 seconds
=== Stratified cross-validation ===
=== Summary ===
Correctly Classified Instances      146      48.6667 %
Kappa statistic                    0.23
Mean absolute error                 0.3422
Root mean squared error             0.585
Relative absolute error             77 %
Root relative squared error         124.0967 %
Total Number of Instances          300

=== Detailed Accuracy By Class ===
          TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Cla
0,460    0,255    0,474    0,460    0,467    0,207    0,603    0,398    A
0,540    0,250    0,519    0,540    0,529    0,287    0,645    0,434    B
0,460    0,265    0,465    0,460    0,462    0,195    0,598    0,394    C
Weighted Avg.  0,487    0,257    0,486    0,487    0,486    0,230    0,615    0,409

=== Confusion Matrix ===
  a  b  c  <-- classified as
46 41 13 | a = A
 6 54 40 | b = B
45  9 46 | c = C
```

The 'Result list' on the bottom left shows the following entries:

- 19:51:33 - functions.LibLINEAR
- 19:51:57 - functions.LibSVM
- 19:52:13 - functions.MultilayerPerceptron

LibSVM

Classifier

Choose **MultilayerPerceptron -L 0.3 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H a**

Test options

☐ Use training set
☐ Supplied test set Set...
☒ Cross-validation Folds **10**
☐ Percentage split % **66**
 More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 19:51:33 - functions.LibLINEAR
- 19:51:57 - functions.LibSVM
- 19:52:13 - functions.MultilayerPerceptron

Classifier output

Time taken to build model: 0.05 seconds

=== Stratified cross-validation ===
 === Summary ===

Correctly Classified Instances	260	86.6667 %
Kappa statistic	0.8	
Mean absolute error	0.0889	
Root mean squared error	0.2981	
Relative absolute error	20 %	
Root relative squared error	63.2456 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0,870	0,060	0,879	0,870	0,874	0,812	0,905	0,808	A
	0,850	0,060	0,876	0,850	0,863	0,796	0,895	0,795	B
	0,880	0,080	0,846	0,880	0,863	0,792	0,900	0,785	C
Weighted Avg.	0,867	0,067	0,867	0,867	0,867	0,800	0,900	0,796	

=== Confusion Matrix ===

a	b	c	<-- classified as
87	4	9	a = A
8	85	7	b = B
4	8	88	c = C

MultilayerPerceptron

Classifier

Choose **MultilayerPerceptron -L 0.3 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H a**

Test options

☐ Use training set
☐ Supplied test set Set...
☒ Cross-validation Folds **10**
☐ Percentage split % **66**
 More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 19:51:33 - functions.LibLINEAR
- 19:51:57 - functions.LibSVM
- 19:52:13 - functions.MultilayerPerceptron

Classifier output

Time taken to build model: 0.22 seconds

=== Stratified cross-validation ===
 === Summary ===

Correctly Classified Instances	209	69.6667 %
Kappa statistic	0.545	
Mean absolute error	0.2696	
Root mean squared error	0.3652	
Relative absolute error	60.6653 %	
Root relative squared error	77.4645 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0,840	0,250	0,627	0,840	0,718	0,559	0,875	0,709	A
	0,720	0,120	0,750	0,720	0,735	0,606	0,882	0,813	B
	0,530	0,085	0,757	0,530	0,624	0,496	0,854	0,785	C
Weighted Avg.	0,697	0,152	0,711	0,697	0,692	0,554	0,870	0,769	

=== Confusion Matrix ===

a	b	c	<-- classified as
84	13	3	a = A
14	72	14	b = B
36	11	53	c = C

We can see that Liblinear has lowest accuracy because in linear classifier we have a single layer which is the output layer. However, we expand with a layer of hidden nodes for A.Neural Network. For example, you can see that there are 72 hidden layer below and the accuracy is much higher than Liblinear.

MultiLayerPerceptron (HiddenLayer =72)

Classifier

Choose **LibSVM** -S 0 -K 2 -D 3 -G 0.0 -R 0.0 -N 0.5 -M 40.0 -C 1.0 -E 0.001 -P 0.1 -model "C:\Program Files\Weka-3-9" -seed 1

Test options

☐ Use training set
☐ Supplied test set Set...
☒ Cross-validation Folds **10**
☐ Percentage split % 66
 More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 19:51:33 - functions.LibLINEAR
- 19:51:57 - functions.LibSVM
- 19:52:13 - functions.MultiLayerPerceptron
- 20:10:20 - functions.MultiLayerPerceptron**

Classifier output

Time taken to build model: 2.36 seconds

=== Stratified cross-validation ===
 === Summary ===

Correctly Classified Instances	293	97.6667 %
Kappa statistic	0.965	
Mean absolute error	0.043	
Root mean squared error	0.1232	
Relative absolute error	9.667 %	
Root relative squared error	26.1266 %	
Total Number of Instances	300	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
Weighted Avg.	0.960	0.010	0.980	0.960	0.970	0.955	0.997	0.995	A
	0.990	0.005	0.990	0.990	0.990	0.985	1.000	0.999	B
	0.980	0.020	0.961	0.980	0.970	0.955	0.998	0.995	C

=== Confusion Matrix ===

	a	b	c	<-- classified as
96	1	3	1	a = A
0	99	1	1	b = B
2	0	98	1	c = C

When we change hidden layers with 72, accuracy increased from 69.6667 to 97.6667.

Assignment#2:

Spiral dataset in Scikit

- Classify the Spiral dataset in Scikit using a Neural Network algorithm
- You need to write code for loading csv dataset files (use the Pandas library)

Solution:

Neural Network Algorithm

```
1 import numpy as np
2 import pandas as pd # library for loading dataset
3 import seaborn as sns #for plotting graphs
4 import matplotlib.pyplot as plt
5 spiral_dataset= pd.read_csv("C:/Program Files/Weka-3-8/data/spiral.csv").values
6 # Assign data to X variable
7 X = spiral_dataset[:,0:2]
8 # Assign data to y variable
9 y = spiral_dataset[:,2]
10 from sklearn.model_selection import train_test_split # Train-Test Split
11 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size= 0.2, random_state=60)
12 from sklearn.preprocessing import StandardScaler #Before making actual predictions,
13 #It is always a good practice to scale the features so that all of them can be uniformly evaluated.
14 scaler = StandardScaler()
15 scaler.fit(X_train)
16 X_train = scaler.transform(X_train)
17 X_test = scaler.transform(X_test)
18 from sklearn.neural_network import MLPClassifier # Model Training and Predictions
19 clf = MLPClassifier(hidden_layer_sizes=(100,100,100), max_iter=1000)
20 clf.fit(X_train, y_train)
21 y_pred = clf.predict(X_test)
22 from sklearn.metrics import classification_report, confusion_matrix #evaluating the the Algorithm
23 print(confusion_matrix(y_test,y_pred))
24 print(classification_report(y_test,y_pred))
```

Output:

```
[[20  0  0]
 [ 0 18  0]
 [ 1  2 19]]
precision    recall  f1-score   support

      0.0      0.95      1.00      0.98         20
      1.0      0.90      1.00      0.95         18
      2.0      1.00      0.86      0.93         22

 accuracy      0.95         60
 macro avg      0.95      0.95      0.95         60
 weighted avg      0.95      0.95      0.95         60
```

Assignment#3:

Diabetes dataset in Scikit

- Classify the Diabetes datasets in Scikit using the Neural Network and Xgboost algorithms
- You need to write code for loading csv dataset files (use the Pandas library)

Solution:

Neural Network Algorithm

```
1 import numpy as np
2 import pandas as pd # Library for loading dataset
3 diabetes_dataset= pd.read_csv("C:/Program Files/Weka-3-8/data/diabetes.csv") #Loading dataset to pandas dataframe
4 diabetes_dataset.head() # information about data set
5 # Assign data from first four columns to X variable
6 X = diabetes_dataset.iloc[:, 0:4]
7 # Assign data from first fifth columns to y variable
8 y = diabetes_dataset.select_dtypes(include=[object])
9 from sklearn.model_selection import train_test_split # Train-Test Split
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
11 from sklearn.preprocessing import StandardScaler # Before making actual predictions
12 #it is always a good practice to scale the features so that all of them can be uniformly evaluated
13 scaler = StandardScaler()
14 scaler.fit(X_train)
15 X_train = scaler.transform(X_train)
16 X_test = scaler.transform(X_test)
17 from sklearn.neural_network import MLPClassifier # Applying Model
18 mlp = MLPClassifier(hidden_layer_sizes=(10,10,10), max_iter=1000)
19 mlp.fit(X_train, y_train.values.ravel())
20 predictions = mlp.predict(X_test)
21 from sklearn.metrics import classification_report, confusion_matrix #evaluating the performance of the model.
22 print(confusion_matrix(y_test,predictions))
23 print(classification_report(y_test,predictions))
```

Output

```
[[92 16]
 [13 33]]
precision    recall  f1-score   support

      NO      0.88      0.85      0.86        108
      YES      0.67      0.72      0.69         46

 accuracy      0.81        154
 macro avg      0.77      0.78      0.78        154
 weighted avg      0.82      0.81      0.81        154
```

Xgboost algorithm

```
1 import numpy as np
2 import pandas as pd # Library for loading dataset
3 diabetes_dataset= pd.read_csv("C:/Program Files/Weka-3-8/data/diabetes.csv") #Loading dataset to pandas dataframe
4 # Assign data from first four columns to X variable
5 X = diabetes_dataset.iloc[:, 0:4]
6 # Assign data from first fifth columns to y variable
7 y = diabetes_dataset.select_dtypes(include=[object])
8 from sklearn.model_selection import train_test_split # Train-Test Split
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
10 from xgboost import XGBClassifier #Applying Model
11 model = XGBClassifier()
12 model.fit(X_train, y_train)
13 print(model)
14 # make predictions for test data
15 predictions = model.predict(X_test)
16 # evaluate predictions
17 from sklearn.metrics import accuracy_score
18 accuracy = accuracy_score(y_test, predictions)
19 print("Accuracy: %.2f%%" % (accuracy * 100.0))
20
```

Output

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0,
              learning_rate=0.1, max_delta_step=0, max_depth=3,
              min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
              nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=None, subsample=1, verbosity=1)
Accuracy: 74.68%
```

Assignment#4:

Diabetes dataset in Keras

- Write code for loading, training and evaluating the Diabetes dataset using a neural network classifier in Keras.

Solution:

What is Keras?

- Keras is a high-level neural network API which is written in Python.
- It is capable of running on top of Tensorflow, CNTK or Theano.
- Keras can be used as a deep learning library. Support Convolutional and Recurrent Neural Networks
- Prototyping with keras is fast and easy
- Runs seamlessly on CPU and GPU

I will build a neural network classification for diabetes dataset classification with keras.

Neural Network Algorithm:

```
1  # Import required libraries
2  import pandas as pd
3  import numpy as np
4  import matplotlib.pyplot as plt
5  import sklearn
6  # Import necessary modules
7  from sklearn.model_selection import train_test_split
8  from sklearn.metrics import mean_squared_error
9  from math import sqrt
10 # Keras specific
11 import keras
12 from tensorflow.keras.models import Sequential
13 from tensorflow.keras.layers import Dense
14 from keras.utils import to_categorical
15 df = pd.read_csv("C:/Program Files/Weka-3-8/data/diabetes.csv") #loading dataset to pandas dataframe
16 #To understand the data better, let's view the dataset details.
17 df.head(3)
18 #we need to check what type of data we have in the dataset
19 print(df.shape)
20 df.describe()
21 #Creating Arrays for the Features and the Response Variable.
22 target_column = ['Diabetes']
23 predictors = list(set(list(df.columns))-set(target_column))
24 df[predictors] = df[predictors]/df[predictors].max()
25 df.describe()
26 #Creating the Training and Test Datasets
27 X = df[predictors].values
28 y = df[target_column].values
29 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=40)
30 print(X_train.shape); print(X_test.shape)
31 # one hot encode outputs
32 y_train = to_categorical(y_train)
33 y_test = to_categorical(y_test)
34 count_classes = y_test.shape[1]
35 print(count_classes)
36 #Define, Compile, and Fit the Keras Classification Model
37 model = Sequential()
38 model.add(Dense(500, activation='relu', input_dim=8))
39 model.add(Dense(100, activation='relu'))
40 model.add(Dense(50, activation='relu'))
41 model.add(Dense(2, activation='softmax'))
42 # Compile the model
43 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```

44 # build the model
45 model.fit(X_train, y_train, batch_size=10, epochs=100)
46 #Predict on the Test Data and Compute Evaluation Metrics;
47 pred_train= model.predict(X_train)
48 scores = model.evaluate(X_train, y_train, verbose=0)
49 print('Accuracy on training data: {}% \n Error on training data: {}'.format(scores[1], 1 - scores[1]))
50 pred_test= model.predict(X_test)
51 scores2 = model.evaluate(X_test, y_test, verbose=0)
52 print('Accuracy on test data: {}% \n Error on test data: {}'.format(scores2[1], 1 - scores2[1]))

```

Outputs:

	0TimesPregnant	PlasmaGlucoseConc	DiastolicBloodPressure	TricepsSkinFoldThickness	SerumInsulin	BMI	DiabetesPedigreeFunction	Age	Diabetes
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1

	0TimesPregnant	PlasmaGlucoseConc	DiastolicBloodPressure	TricepsSkinFoldThickness	SerumInsulin	BMI	DiabetesPedigreeFunction	Age	Diabetes
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

#Creating the Training and Test Datasets

```

X = df[predictors].values
y = df[target_column].values

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=40)
print(X_train.shape); print(X_test.shape)

```

(537, 8)

(231, 8)

one hot encode outputs

```

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

```

```

count_classes = y_test.shape[1]
print(count_classes)

```

2

537/537 [=====] - 1s 2ms/sample - loss: 0.2322 - accuracy: 0.9069

Epoch 93/100

537/537 [=====] - 1s 1ms/sample - loss: 0.2184 - accuracy: 0.9069

Epoch 94/100

537/537 [=====] - 1s 2ms/sample - loss: 0.2066 - accuracy: 0.9181

Epoch 95/100

537/537 [=====] - 0s 917us/sample - loss: 0.2019 - accuracy: 0.9199

Epoch 96/100

537/537 [=====] - 0s 836us/sample - loss: 0.2040 - accuracy: 0.9255

Epoch 97/100

537/537 [=====] - 1s 1ms/sample - loss: 0.2021 - accuracy: 0.9218

Epoch 98/100

537/537 [=====] - 0s 867us/sample - loss: 0.2023 - accuracy: 0.9199

Epoch 99/100

537/537 [=====] - 0s 650us/sample - loss: 0.2042 - accuracy: 0.9106 - loss: 0.1768 - accuracy: 0.9106

Epoch 100/100

537/537 [=====] - 1s 982us/sample - loss: 0.2074 - accuracy: 0.9125

Accuracy on training data: 0.9199255108833313

Error on training data: 0.0800744891166687

Accuracy on test data: 0.701298713684082%

Error on test data: 0.29870128631591797

Assignment#5:

MNSIT dataset in Keras

- Write code for loading, training and evaluating the MNIST dataset using a Linear, Neural Network and a ConvNet classifier in Keras.

Solution:

We're going to tackle a classic machine learning problem: MNIST handwritten digit classification. It's simple: given an image, classify it as a digit.



Sample images from the MNIST dataset

Each image in the MNIST dataset is 28x28 and contains a centered, grayscale digit. We'll flatten each 28x28 into a 784 dimensional vector, which we'll use as input to our neural network. Our output will be one of 10 possible classes: one for each digit.

Linear

```
In [1]: import numpy as np
import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from keras.utils import to_categorical

Using TensorFlow backend.
```

```
In [2]: train_images = mnist.train_images()
train_labels = mnist.train_labels()
test_images = mnist.test_images()
test_labels = mnist.test_labels()
```

```
In [3]: # Normalize the images.
train_images = (train_images / 255) - 0.5
test_images = (test_images / 255) - 0.5
```

```
In [4]: # Flatten the images.
train_images = train_images.reshape((-1, 784))
test_images = test_images.reshape((-1, 784))
```

```
In [5]: # Build the model.
model = Sequential([
    Dense(10, activation='softmax'),
])
```

```
In [6]: # Compile the model.
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
```

```
In [7]: # Train the model.
model.fit(
    train_images,
    to_categorical(train_labels),
    epochs=5,
    batch_size=32,
)
```

Train on 60000 samples

Epoch 1/5

60000/60000 [=====] - 188s 3ms/sample - loss: 0.4924 - accuracy: 0.8652

Epoch 2/5

60000/60000 [=====] - 188s 3ms/sample - loss: 0.3323 - accuracy: 0.9037

Epoch 3/5

60000/60000 [=====] - 187s 3ms/sample - loss: 0.3113 - accuracy: 0.9098

Epoch 4/5

60000/60000 [=====] - 187s 3ms/sample - loss: 0.3010 - accuracy: 0.9137

Epoch 5/5

60000/60000 [=====] - 187s 3ms/sample - loss: 0.2937 - accuracy: 0.9172

Out[7]: <tensorflow.python.keras.callbacks.History at 0x209e7a6ee48>

Neural Network

```
In [1]: import numpy as np
import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from keras.utils import to_categorical
```

Using TensorFlow backend.

```
In [2]: train_images = mnist.train_images()
train_labels = mnist.train_labels()
test_images = mnist.test_images()
test_labels = mnist.test_labels()
```

```
In [3]: # Normalize the images.
train_images = (train_images / 255) - 0.5
test_images = (test_images / 255) - 0.5
```

```
In [4]: # Flatten the images.
train_images = train_images.reshape((-1, 784))
test_images = test_images.reshape((-1, 784))
```

```
In [5]: # Build the model.
model = Sequential([
    Dense(64, activation='relu', input_shape=(784,)),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax'),
])
```

```
In [6]: # Compile the model.
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)
```

```
In [7]: # Train the model.
model.fit(
    train_images,
    to_categorical(train_labels),
    epochs=5,
    batch_size=32,
)
```

Train on 60000 samples

Epoch 1/5

60000/60000 [=====] - 189s 3ms/sample - loss: 0.3666 - accuracy: 0.8901

Epoch 2/5

60000/60000 [=====] - 172s 3ms/sample - loss: 0.1911 - accuracy: 0.9417

Epoch 3/5

60000/60000 [=====] - 184s 3ms/sample - loss: 0.1493 - accuracy: 0.9538

Epoch 4/5

60000/60000 [=====] - 178s 3ms/sample - loss: 0.1262 - accuracy: 0.9617

Epoch 5/5

60000/60000 [=====] - 166s 3ms/sample - loss: 0.1112 - accuracy: 0.9655

Out[7]: <tensorflow.python.keras.callbacks.History at 0x28a02c04048>

In [8]: *# Evaluate the model.*

```
model.evaluate(
    test_images,
    to_categorical(test_labels)
)
```

3s 330us/sample - loss: 0.0723 - accuracy: 0.9620

ConvNet

In [1]: *#Loading required libraries*

```
import numpy as np
import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten
from keras.utils import to_categorical
```

Using TensorFlow backend.

In [2]: `train_images = mnist.train_images()`
`train_labels = mnist.train_labels()`
`test_images = mnist.test_images()`
`test_labels = mnist.test_labels()`

In [3]: *# Normalize the images.*

```
train_images = (train_images / 255) - 0.5
test_images = (test_images / 255) - 0.5
```

In [4]: *# Reshape the images.*

```
train_images = np.expand_dims(train_images, axis=3)
test_images = np.expand_dims(test_images, axis=3)
num_filters = 8
filter_size = 3
pool_size = 2
```

In [5]: *# Build the model.*

```
model = Sequential([
    Conv2D(num_filters, filter_size, input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=pool_size),
    Flatten(),
    Dense(10, activation='softmax'),
])
```

In [6]: *# Compile the model.*

```
model.compile(
    'adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)
```

In [7]: *# Train the model.*

```
model.fit(
    train_images,
    to_categorical(train_labels),
    epochs=3,
    validation_data=(test_images, to_categorical(test_labels)),
)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/3

60000/60000 [=====] - 210s 4ms/sample - loss: 0.3673 - accuracy: 0.8932 - val_loss: 0.2291 - val_accuracy: 0.9335

Epoch 2/3

60000/60000 [=====] - 205s 3ms/sample - loss: 0.2001 - accuracy: 0.9418 - val_loss: 0.1622 - val_accuracy: 0.9535

Epoch 3/3

60000/60000 [=====] - 197s 3ms/sample - loss: 0.1453 - accuracy: 0.9585 - val_loss: 0.1395 - val_accuracy: 0.9606

Out[7]: <tensorflow.python.keras.callbacks.History at 0x1fe7c369f08>

Assignment#6:

Banknote dataset in Web ML Experimenter

- Download the Banknote dataset from the Datasets page
- Upload the dataset in the Web ML Experimenter
- Try classifying the Banknote dataset using different classifiers. Test with different hyperparameter settings.
- Which classifier had the highest accuracy?

Solution:

k-NN

Upload Dataset ?

Dosya Seç Loaded dataset: [banknote.csv](#) [Try Iris dataset](#)

Select Algorithm


☒ k-NN ☐ Linear ☐ Neural Network ☐ Decision Tree ☐ Random Forest ☐ SVM ☐ Naive Bayes

Set Hyperparameters ?

Neighbors: Distance function:

Data Pre-processing ?

Run Experiment ?

 ☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

N=3

Result ?

File: banknote.csv
Number of examples: 1372
Number of attributes: 4
Number of classes: 2

Accuracy by fold:

1	2	3	4	5
100.00%	99.27%	100.00%	100.00%	100.00%

Total accuracy: 99.85% (1370/1372 correctly classified)

Evaluation time: 1.720 sec

Confusion Matrix:

	[0]	[1]	
[0]	760	2	→ 0
[1]	0	610	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	1.000	0.997	0.999	→ 0
[1]	0.997	1.000	0.998	→ 1
Avg:	0.998	0.999	0.999	

N=7

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
100.00%	99.27%	99.64%	100.00%	100.00%

Total accuracy: 99.78% (1369/1372 correctly classified)

Evaluation time: 1.482 sec

Confusion Matrix:

	[0]	[1]	
[0]	759	3	→ 0
[1]	0	610	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	1.000	0.996	0.998	→ 0
[1]	0.995	1.000	0.998	→ 1
Avg:	0.998	0.998	0.998	

Linear

Upload Dataset ?

Dosya Seç banknote.csv

Loaded dataset: banknote.csv

[Try Iris dataset](#)

Select Algorithm

☐ k-NN ☒ Linear ☐ Neural Network ☐ Decision Tree ☐ Random Forest ☐ SVM ☐ Naïve Bayes

Set Hyperparameters ?

Training iterations: Learning rate: Batch size:
L2 regularization: Momentum:

Data Pre-processing ?

Run Experiment ?



☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Training iteration : 1000 – Batch Size: 0

Result ?

File: banknote.csv
Number of examples: 1372
Number of attributes: 4
Number of classes: 2

Accuracy by fold:

1	2	3	4	5
97.45%	99.64%	97.81%	97.81%	99.64%

Total accuracy: 98.47% (1351/1372 correctly classified)

Evaluation time: 3.469 sec

Confusion Matrix:

	[0]	[1]	
[0]	744	18	→ 0
[1]	3	607	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.996	0.976	0.986	→ 0
[1]	0.971	0.995	0.983	→ 1
Avg:	0.984	0.986	0.985	

Training iteration : 1000 – Batch Size: 25

Result ?

File: banknote.csv
Number of examples: 1372
Number of attributes: 4
Number of classes: 2

Accuracy by fold:

1	2	3	4	5
97.45%	98.91%	97.81%	97.81%	99.64%

Total accuracy: 98.32% (1349/1372 correctly classified)

Evaluation time: 1.742 sec

Confusion Matrix:

	[0]	[1]	
[0]	742	20	→ 0
[1]	3	607	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.996	0.974	0.985	→ 0
[1]	0.968	0.995	0.981	→ 1
Avg:	0.982	0.984	0.983	

Training iteration : 500 – Batch Size: 25

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
97.08%	98.91%	97.81%	97.81%	98.91%

Total accuracy: 98.10% (1346/1372 correctly classified)

Evaluation time: 0.890 sec

Confusion Matrix:

	[0]	[1]	
[0]	740	22	→ 0
[1]	4	606	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.995	0.971	0.983	→ 0
[1]	0.965	0.993	0.979	→ 1
Avg:	0.980	0.982	0.981	

Neural Network

Upload Dataset ?

Dosya Seç banknote.csv

Loaded dataset: banknote.csv

[Try Iris dataset](#)

Select Algorithm

☐ k-NN ☐ Linear ☒ Neural Network ☐ Decision Tree ☐ Random Forest ☐ SVM ☐ Naïve Bayes

▼ Set Hyperparameters ?

Training iterations: 500 Hidden layers: 32 Learning rate: 0.8 Batch size: 0
L2 regularization: 0 Momentum: 0.1 Dropout: 0.2

► Data Pre-processing ?

Run Experiment ?



☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Training iteration : 500 – Batch Size: 0 – Hidden Layers: 32

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1

2

3

4

5

100.00%	100.00%	100.00%	100.00%	100.00%
---------	---------	---------	---------	---------

Total accuracy: 100.00% (1372/1372 correctly classified)

Evaluation time: 18.739 sec

Confusion Matrix:

	[0]	[1]	
[0]	762	0	→ 0
[1]	0	610	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	1.000	1.000	1.000	→ 0
[1]	1.000	1.000	1.000	→ 1
Avg:	1.000	1.000	1.000	

Training iteration : 500 – Batch Size: 25 – Hidden Layers: 32

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1

2

3

4

5

98.91%

100.00%

98.54%

100.00%

100.00%

Total accuracy: 99.49% (1365/1372 correctly classified)

Evaluation time: 7.867 sec

Confusion Matrix:

[0]

[1]

[0]

755

7

→ 0

[1]

0

610

→ 1

Metrics by category:

Precision

Recall

F1 score

[0]

1.000

0.991

0.995

→ 0

[1]

0.989

1.000

0.994

→ 1

Avg:

0.994

0.995

0.995

Training iteration : 100 – Batch Size: 25 – Hidden Layers: 32

Result ?				
File: banknote.csv				
Number of examples: 1372				
Number of attributes: 4				
Number of classes: 2				
Accuracy by fold:				
1	2	3	4	5
98.18%	99.64%	97.81%	98.54%	98.91%
Total accuracy: 98.62% (1353/1372 correctly classified)				
Evaluation time: 1.654 sec				
Confusion Matrix:				
	[0]	[1]		
[0]	743	19	→ 0	
[1]	0	610	→ 1	
Metrics by category:				
	Precision	Recall	F1 score	
[0]	1.000	0.975	0.987	→ 0
[1]	0.970	1.000	0.985	→ 1
Avg:	0.985	0.988	0.986	

Training iteration : 100 – Batch Size: 25 – Hidden Layers: 16

Result ?				
File: banknote.csv				
Number of examples: 1372				
Number of attributes: 4				
Number of classes: 2				
Accuracy by fold:				
1	2	3	4	5
98.18%	99.27%	98.18%	98.18%	98.91%
Total accuracy: 98.54% (1352/1372 correctly classified)				
Evaluation time: 0.882 sec				
Confusion Matrix:				
	[0]	[1]		
[0]	745	17	→ 0	
[1]	3	607	→ 1	
Metrics by category:				
	Precision	Recall	F1 score	
[0]	0.996	0.978	0.987	→ 0
[1]	0.973	0.995	0.984	→ 1
Avg:	0.984	0.986	0.985	

Decision Tree

Upload Dataset ?

Dosya Seç banknote.csv

Loaded dataset: banknote.csv

[Try Iris dataset](#)

Select Algorithm

☐ k-NN ☐ Linear ☐ Neural Network ☒ Decision Tree ☐ Random Forest ☐ SVM ☐ Naive Bayes

Set Hyperparameters ?

Max depth:

7

Min samples split:

5

Data Pre-processing ?

Run Experiment ?



☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Max Depth: 7 – Min samples split: 5

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
98.18%	100.00%	100.00%	100.00%	100.00%

Total accuracy: 99.64% (1367/1372 correctly classified)

Evaluation time: 19.321 sec

Confusion Matrix:

	[0]	[1]	
[0]	760	2	→ 0
[1]	3	607	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.996	0.997	0.997	→ 0
[1]	0.997	0.995	0.996	→ 1
Avg:	0.996	0.996	0.996	

Max Depth: 3 – Min samples split: 5

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
95.26%	93.80%	93.80%	91.97%	96.38%

Total accuracy: 94.24% (1293/1372 correctly classified)

Evaluation time: 14.734 sec

Confusion Matrix:

	[0]	[1]	
[0]	728	34	→ 0
[1]	45	565	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.942	0.955	0.949	→ 0
[1]	0.943	0.926	0.935	→ 1
Avg:	0.943	0.941	0.942	

Random Forest

Upload Dataset ?

Dosya Seç

banknote.csv

Loaded dataset: banknote.csv

Try Iris dataset

Select Algorithm

☐ k-NN ☐ Linear ☐ Neural Network ☐ Decision Tree ☒ Random Forest ☐ SVM ☐ Naive Bayes

▼ Set Hyperparameters ?

No estimators:

15

Max depth:

7

Min samples split:

5

Subset sample size:

0.9

► Data Pre-processing ?

Run Experiment ?

☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Max Depth: 7 – Min samples split: 5

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1

2

3

4

5

99.27%

100.00%

100.00%

100.00%

99.64%

Total accuracy: 99.78% (1369/1372 correctly classified)

Evaluation time: 14.858 sec

Confusion Matrix:

[0]

[1]

[0]

760

2

→ 0

[1]

1

609

→ 1

Metrics by category:

Precision

Recall

F1 score

[0]

0.999

0.997

0.998

→ 0

[1]

0.997

0.998

0.998

→ 1

Avg:

0.998

0.998

0.998

Max Depth: 3 – Min samples split: 5

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
94.89%	92.70%	93.43%	92.70%	96.74%

Total accuracy: 94.10% (1291/1372 correctly classified)

Evaluation time: 11.667 sec

Confusion Matrix:

	[0]	[1]	
[0]	705	57	→ 0
[1]	24	586	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.967	0.925	0.946	→ 0
[1]	0.911	0.961	0.935	→ 1
Avg:	0.939	0.943	0.941	

SVM

Upload Dataset ?

Dosya Seç banknote.csv

Loaded dataset: banknote.csv

Try Iris dataset

Select Algorithm

☐ k-NN ☐ Linear ☐ Neural Network ☐ Decision Tree ☐ Random Forest ☒ SVM ☐ Naïve Bayes

Set Hyperparameters ?

Gamma: 5

Data Pre-processing ?

Run Experiment ?



☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Gamma = 5

Result ?

File: banknote.csv

Number of examples: 1372

Number of attributes: 4

Number of classes: 2

Accuracy by fold:

1	2	3	4	5
98.54%	100.00%	99.64%	99.64%	99.64%

Total accuracy: 99.49% (1365/1372 correctly classified)

Evaluation time: 0.944 sec

Confusion Matrix:

	[0]	[1]	
[0]	757	5	→ 0
[1]	2	608	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.997	0.993	0.995	→ 0
[1]	0.992	0.997	0.994	→ 1
Avg:	0.995	0.995	0.995	

Gamma = 7

Result ?

File: banknote.csv
Number of examples: 1372
Number of attributes: 4
Number of classes: 2

Accuracy by fold:

1	2	3	4	5
99.27%	100.00%	100.00%	100.00%	100.00%

Total accuracy: 99.85% (1370/1372 correctly classified)

Evaluation time: 0.940 sec

Confusion Matrix:

	[0]	[1]	
[0]	762	0	→ 0
[1]	2	608	→ 1

Metrics by category:

	Precision	Recall	F1 score	
[0]	0.997	1.000	0.999	→ 0
[1]	1.000	0.997	0.998	→ 1
Avg:	0.999	0.998	0.999	

Naive Bayes

Upload Dataset ?

Dosya Seç

banknote.csv

Loaded dataset: banknote.csv

Try Iris dataset

Select Algorithm

☐ k-NN ☐ Linear ☐ Neural Network ☐ Decision Tree ☐ Random Forest ☐ SVM ☒ Naive Bayes

▼ Set Hyperparameters ?

► Data Pre-processing ?

Run Experiment ?

☐ Use training data ☐ Train-test split (80%/20%) ☒ 5-fold Cross-Validation

Result ?				
File: banknote.csv				
Number of examples: 1372				
Number of attributes: 4				
Number of classes: 2				
Accuracy by fold:				
1	2	3	4	5
86.86%	81.75%	80.66%	83.21%	90.22%
Total accuracy: 84.55% (1160/1372 correctly classified)				
Evaluation time: 0.051 sec				
Confusion Matrix:				
	[0]	[1]		
[0]	662	100	→ 0	
[1]	112	498	→ 1	
Metrics by category:				
	Precision	Recall	F1 score	
[0]	0.855	0.869	0.862	→ 0
[1]	0.833	0.816	0.825	→ 1
Avg:	0.844	0.843	0.843	

According to results, we can easily see that Naive Bayes Classification has the worst accuracy. Besides, we can see that Neural Network Classification has the highest accuracy because it is the only one which we get 100% percent of accuracy.

Assignment#7:

Diabetes dataset in R

- Classify the diabetes dataset in R using Neural Networks, SVM and RandomForest.
- Split the dataset into 80% training and 20% testing, and evaluate accuracy on the test dataset
- Which classifier had the highest accuracy?

Solution:

Support Vector Machine

```
1 # Read the Data
2 dataset <- read.csv("C:\\Program Files\\weka-3-8\\data\\diabetes.csv", sep = ',', header = FALSE)
3 str(dataset)
4 head(dataset)
5 #80% training and 20% testing
6 intrain <- createDataPartition(y = dataset$V9, p= 0.8, list = FALSE)
7 training <- dataset[intrain,]
8 testing <- dataset[-intrain,]
9 dim(training);
10 dim(testing);
11 anyNA(dataset)
12 summary(dataset)
13 training[["V9"]] = factor(training[["V9"]])
14 trctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 3)
15 #Model
16 svm_Linear <- train(V9 ~., data = training, method = "svmLinear",
17                     trControl=trctrl,
18                     preProcess = c("center", "scale"),
19                     tuneLength = 10)
20 svm_Linear
21 test_pred <- predict(svm_Linear, newdata = testing)
22 test_pred
23 #Evaluation
24 confusionMatrix(table(test_pred, testing$V9))
25
```

Output:

```
> str(dataset)
'data.frame': 769 obs. of 9 variables:
 $ V1: Factor w/ 18 levels "0","1","10","11",...: 18 14 2 16 2 1 13 11 3 10 ...
 $ V2: Factor w/ 137 levels "0","100","101",...: 137 50 122 85 126 39 18 115 17 97 ...
 $ V3: Factor w/ 48 levels "0","100","102",...: 48 31 28 26 28 13 32 17 1 30 ...
 $ V4: Factor w/ 52 levels "0","10","11",...: 52 27 21 1 15 27 1 24 1 37 ...
 $ V5: Factor w/ 187 levels "0","100","105",...: 187 1 1 1 183 41 1 178 1 141 ...
 $ V6: Factor w/ 249 levels "0.0","18.2","18.4",...: 249 124 63 31 78 210 54 104 141 100 ...
 $ V7: Factor w/ 518 levels "0.078","0.084",...: 518 351 197 369 54 515 81 119 24 45 ...
 $ V8: Factor w/ 53 levels "21","22","23",...: 53 30 11 12 1 13 10 6 9 33 ...
 $ V9: Factor w/ 3 levels "Diabetes","No",...: 1 3 2 3 2 3 2 3 2 3 ...
> head(dataset)
   V1      V2      V3      V4      V5      V6
1 NoTimesPregnant PlasmaGlucoseConc DiastolicBloodPressure TricepsSkinFoldThickness SerumInsulin BMI
2 6 148 72 35 0 33.6
3 1 85 66 29 0 26.6
4 8 183 64 0 0 23.3
5 1 89 66 23 94 28.1
6 0 137 40 35 168 43.1
   V7      V8      V9
1 DiabetesPedigreeFunction Age Diabetes
2 0.627 50 YES
3 0.351 31 NO
4 0.672 32 YES
5 0.167 21 NO
6 2.288 33 YES
```

```
> dim(training);
[1] 616 9
> dim(testing);
[1] 153 9
> anyNA(dataset)
[1] FALSE
> summary(dataset)
   V1      V2      V3      V4      V5      V6      V7      V8
1 :135 100 : 17 70 : 57 0 :227 0 :374 32.0 : 13 0.254 : 6 22 : 72
0 :111 99 : 17 74 : 52 32 : 31 105 : 11 31.2 : 12 0.258 : 6 21 : 63
2 :103 106 : 14 68 : 45 30 : 27 130 : 9 31.6 : 12 0.207 : 5 25 : 48
3 : 75 111 : 14 78 : 45 27 : 23 140 : 9 0.0 : 11 0.238 : 5 24 : 46
4 : 68 125 : 14 72 : 44 23 : 22 120 : 8 32.4 : 10 0.259 : 5 23 : 38
5 : 57 129 : 14 64 : 43 18 : 20 100 : 7 33.3 : 10 0.261 : 5 28 : 35
(other):220 (other):679 (other):483 (other):419 (other):351 (other):701 (other):737 (other):467
   V9
Diabetes: 1
NO :500
YES :268
```

```
> svm_Linear
Support Vector Machines with Linear Kernel

616 samples
8 predictor
3 classes: 'Diabetes', 'NO', 'YES'

Pre-processing: centered (1254), scaled (1254)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 554, 554, 554, 554, 553, 555, ...
Resampling results:

Accuracy      Kappa
0.5691608     0.1112582
```


Confusion Matrix and Statistics

```
test_pred Diabetes NO YES
Diabetes    0  0  0
NO          0 79 32
YES         0 21 21
```

Overall Statistics

```
Accuracy : 0.6536
95% CI : (0.5725, 0.7286)
No Information Rate : 0.6536
P-Value [Acc > NIR] : 0.5373
```

Kappa : 0.1958

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: Diabetes	Class: NO	Class: YES
Sensitivity	NA	0.7900	0.3962
Specificity	1	0.3962	0.7900
Pos Pred Value	NA	0.7117	0.5000
Neg Pred Value	NA	0.5000	0.7117
Prevalence	0	0.6536	0.3464
Detection Rate	0	0.5163	0.1373
Detection Prevalence	0	0.7255	0.2745
Balanced Accuracy	NA	0.5931	0.5931

Random Forest

```
1 # Read the Data
2 dataset <- read.csv("C:\\Program Files\\weka-3-8\\data\\diabetes.csv", sep = ',', header = TRUE)
3 # Test-Training Split = 80 : 20
4 set.seed(100)
5 train <- sample(nrow(dataset), 0.8*nrow(dataset), replace = FALSE)
6 Trainset <- dataset[train,]
7 validset <- dataset[-train,]
8 # Create a Random Forest model with default parameters
9 model1 <- randomForest(Diabetes ~ ., data = Trainset, importance = TRUE)
10 model1
11 # Predicting on train set
12 predTrain <- predict(model1, Trainset, type = "class")
13 # Checking classification accuracy
14 table(predTrain, Trainset$Diabetes)
15 # Predicting on Validation set
16 predValid <- predict(model1, validset, type = "class")
17 table(predValid, validset$Diabetes)
18 # Checking classification accuracy
19 mean(predValid == validset$Diabetes)
```

Output:

> model1

Call:

```
randomForest(formula = Diabetes ~ ., data = Trainset, importance = TRUE)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 2

OOB estimate of error rate: 23.13%

Confusion matrix:

	NO	YES	class.error
NO	348	54	0.1343284
YES	88	124	0.4150943

```
> # Predicting on train set
> predTrain <- predict(model1, Trainset, type = "class")
> # Checking classification accuracy
> table(predTrain, Trainset$Diabetes)

predTrain NO YES
NO 402 0
YES 0 212
> # Predicting on Validation set
> predValid <- predict(model1, validset, type = "class")
> # Checking classification accuracy
> table(predValid, validset$Diabetes)

predValid NO YES
NO 86 25
YES 12 31
> mean(predValid == validset$Diabetes)
[1] 0.7597403
```

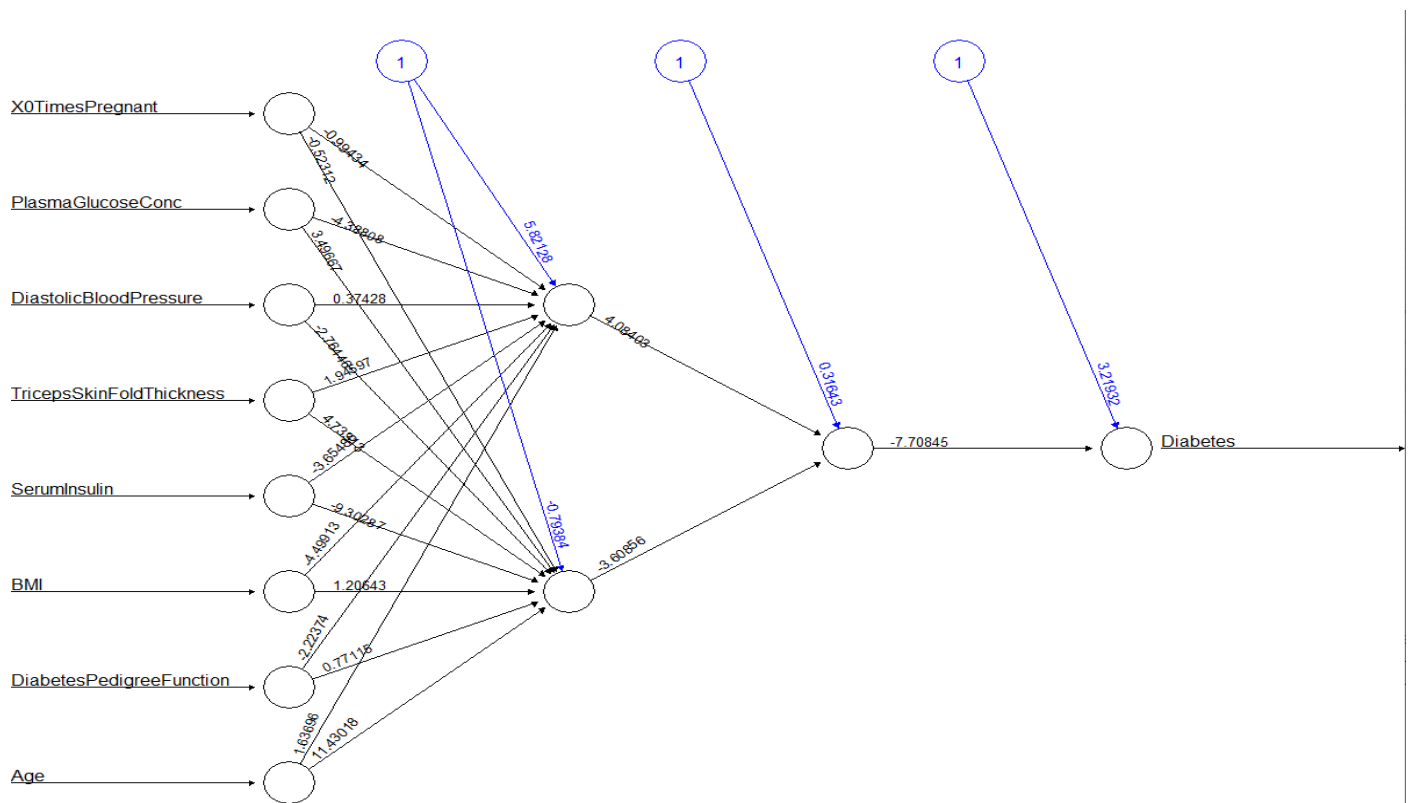
Neural Network

```

1 # Read the Data
2 diabetes <- read.csv("C:\\Program Files\\Weka-3-8\\data\\diabetes-.csv", sep = ',', header = TRUE)
3 attach(diabetes)
4
5 #Scaled Normalization
6 scaleddata<-scale(diabetes)
7
8 #Max-Min Normalization
9 normalize <- function(x) {
10   return ((x - min(x)) / (max(x) - min(x)))
11 }
12 maxmindf <- as.data.frame(lapply(diabetes, normalize))
13
14 # Training and Test Data %80-%20
15 trainset <- maxmindf[1:614, ]
16 testset <- maxmindf[615:768, ]
17
18 #Training a Neural Network Model using neuralnet
19 library(neuralnet)
20 nn <- neuralnet(Diabetes ~ X0TimesPregnant + PlasmaGlucoseConc + DiastolicBloodPressure +
21   TricepsSkinFoldThickness + SerumInsulin + BMI + DiabetesPedigreeFunction +Age, data=trainset, hidden=c(2,1), linear.output=FALSE, threshold=0.01)
22 nn$result.matrix
23 plot(nn)
24
25 #Testing The Accuracy Of The Model
26 temp_test <- subset(testset, select = c("X0TimesPregnant", "PlasmaGlucoseConc", "DiastolicBloodPressure", "TricepsSkinFoldThickness",
27   "SerumInsulin", "BMI", "DiabetesPedigreeFunction", "Age"))
28 head(temp_test)
29 nn.results <- compute(nn, temp_test)
30 results <- data.frame(actual = testset$Diabetes, prediction = nn.results$net.result)
31
32 #Confusion Matrix
33 roundedresults<-sapply(results,round,digits=0)
34 roundedresultsdf=data.frame(roundedresults)
35 attach(roundedresultsdf)
36 table(actual,prediction)

```

Output:



Error: 43.200533 Steps: 4729

```

> head(temp_test)
  X0TimesPregnant PlasmaGlucoseConc DiastolicBloodPressure TricepsSkinFoldThickness SerumInsulin BMI
615      0.6470588      0.6934673      0.6065574      0.2626263      0.1702128 0.5380030
616      0.1764706      0.5326633      0.5901639      0.0000000      0.0000000 0.3845007
617      0.3529412      0.5879397      0.7868852      0.0000000      0.0000000 0.4277198
618      0.1176471      0.3417085      0.5081967      0.1313131      0.0177305 0.2995529
619      0.5294118      0.5628141      0.6721311      0.2424242      0.0000000 0.4202683
620      0.0000000      0.5979899      0.0000000      0.0000000      0.0000000 0.4828614

  DiabetesPedigreeFunction Age
615      0.20452605 0.48333333
616      0.05508113 0.10000000
617      0.03373185 0.15000000
618      0.07643040 0.03333333
619      0.51409052 0.48333333
620      0.02690009 0.05000000

> nn$result.matrix
      [,1]
error      4.320053e+01
reached.threshold 9.937623e-03
steps      4.729000e+03
Intercept.to.1layhid1 5.821279e+00
X0TimesPregnant.to.1layhid1 -9.943362e-01
PlasmaGlucoseConc.to.1layhid1 -4.388082e+00
DiastolicBloodPressure.to.1layhid1 3.742786e-01
TricepsSkinFoldThickness.to.1layhid1 1.945970e+00
SerumInsulin.to.1layhid1 -3.654856e+00
BMI.to.1layhid1 -4.499135e+00
DiabetesPedigreeFunction.to.1layhid1 -2.223742e+00
Age.to.1layhid1 1.636956e+00
Intercept.to.1layhid2 -7.938379e-01
X0TimesPregnant.to.1layhid2 -5.231193e-01
PlasmaGlucoseConc.to.1layhid2 3.496672e+00
DiastolicBloodPressure.to.1layhid2 -2.764459e+00
TricepsSkinFoldThickness.to.1layhid2 4.733127e+00
SerumInsulin.to.1layhid2 -9.302873e+00
BMI.to.1layhid2 1.206430e+00
DiabetesPedigreeFunction.to.1layhid2 7.711523e-01
Age.to.1layhid2 1.143018e+01
Intercept.to.2layhid1 3.164293e-01
1layhid1.to.2layhid1 4.084031e+00
1layhid2.to.2layhid1 -3.608559e+00
Intercept.to.Diabetes 3.219325e+00
2layhid1.to.Diabetes -7.708448e+00

$net.result
      [,1]
615 0.77665766
616 0.06212715
617 0.09912713
618 0.02287310
619 0.37935852
620 0.34133622
621 0.19079303
622 0.24055323
623 0.93723970
624 0.06269083
625 0.03862114
626 0.16018191
627 0.04205176
628 0.10494110
629 0.39292667
630 0.03966441
631 0.38773735
632 0.08916785
633 0.05835261
634 0.02382234
635 0.09581610

> table(actual,prediction)
      prediction
actual 0 1
0 86 13
1 21 34

> (86+34)/(86+34+21+13)
[1] 0.7792208

```

As you see that neural network classification has the highest accuracy.

Assignment#8:

Pre-trained models in Keras

- Use the pre-trained models VGG16 and VGG19 in Keras to classify images.
- Test on other images than the examples. Are they classified correctly?

Solution:

Convolutional neural networks are now capable of outperforming humans on some computer vision tasks, such as classifying images. That is, given a photograph of an object, answer the question as to which of 1,000 specific objects the photograph shows. A competition-winning model for this task is the VGG model by researchers at Oxford. What is important about this model, besides its capability of classifying objects in photographs, is that the model weights are freely available and can be loaded and used in your own models and applications.

VGG16

I used water bottle, car ,signboard images for testing VGG16 algorithm.

```
In [1]: from tensorflow.keras.applications.vgg16 import VGG16
model = VGG16()
```

```
In [2]: print(model.summary())
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 224, 224, 3)]	0

block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928

block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584

block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080

block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080

block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0

block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
=====		
Total params: 138,357,544		
Trainable params: 138,357,544		
Non-trainable params: 0		

```
In [3]: from keras.preprocessing.image import load_img
# load an image from file
image = load_img('araba.jpg', target_size=(224, 224))
```

Using TensorFlow backend.

```
In [4]: from keras.preprocessing.image import img_to_array
# convert the image pixels to a numpy array
image = img_to_array(image)
```

```
In [5]: # reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
```

```
In [6]: from keras.applications.vgg16 import preprocess_input
# prepare the image for the VGG model
image = preprocess_input(image)
```

```
In [7]: # predict the probability across all output classes
prob = model.predict(image)
```

```
In [8]: from keras.applications.vgg16 import decode_predictions
# convert the probabilities to class labels
label = decode_predictions(prob)
# retrieve the most likely result, e.g. highest probability
label = label[0][0]
# print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))
```

OUTPUTS:



racer (61.10%)



water_bottle (99.59%)



street_sign (90.77%)

VGG19

I used water bottle, car, signboard, monkey, dog, cat, jeep images for testing VGG19 algorithm.

```
In [1]: from tensorflow.keras.applications.vgg19 import VGG19
        model = VGG19()
```

```
In [2]: print(model.summary())
```

Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
Total params: 143,667,240		
Trainable params: 143,667,240		
Non-trainable params: 0		


```
In [3]: from keras.preprocessing.image import load_img
# load an image from file
image = load_img('signboard.jpg', target_size=(224, 224))

Using TensorFlow backend.
```

```
In [4]: from keras.preprocessing.image import img_to_array
# convert the image pixels to a numpy array
image = img_to_array(image)
```

```
In [5]: # reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
```

```
In [6]: from keras.applications.vgg16 import preprocess_input
# prepare the image for the VGG model
image = preprocess_input(image)
```

```
In [7]: # predict the probability across all output classes
prob = model.predict(image)
```

```
In [8]: from keras.applications.vgg16 import decode_predictions
# convert the probabilities to class labels
label = decode_predictions(prob)
# retrieve the most likely result, e.g. highest probability
label = label[0][0]
# print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))
```



street_sign (91.03%)



sports_car (52.97%)



water_bottle (98.88%)



tabby (51.24%)

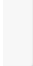


Labrador_retriever (36.40%)



patas (94.09%)



 jeep (29.14%)

As you see, they are correctly classified, the algorithm even can understand type of cars and different species of animals.